


Reinforcement Learning (Single Agent Q-learning)

ECE 277

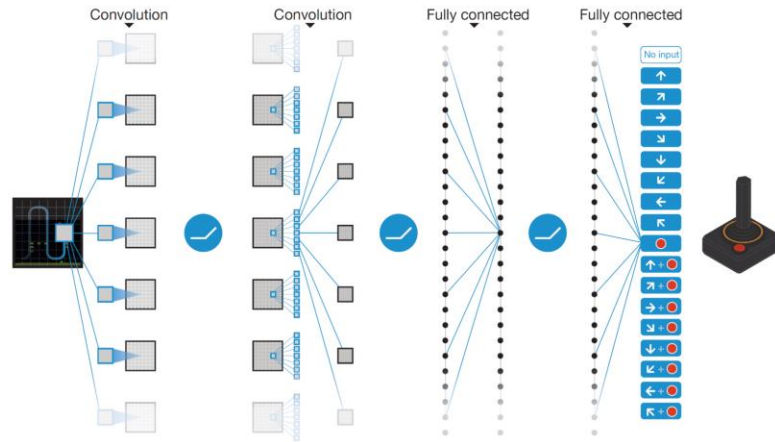
Cheolhong An

Deep Reinforcement learning: RL + Deep Learning



Learned to play 49 games for the Atari 2600 game console,
without labels or human input, from self-play and the score alone

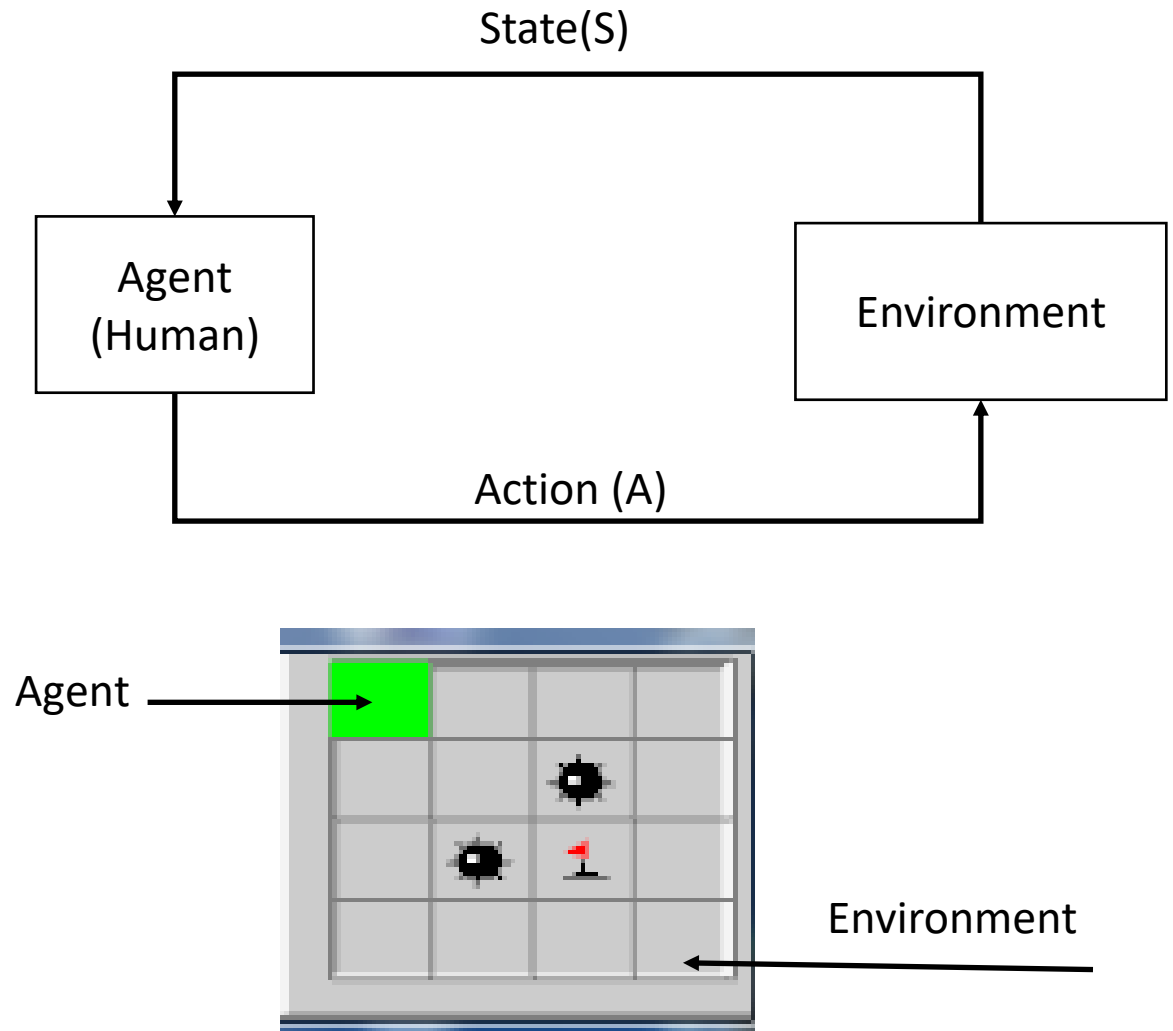
mapping raw
screen pixels



to predictions
of final score
for each of 18
joystick actions


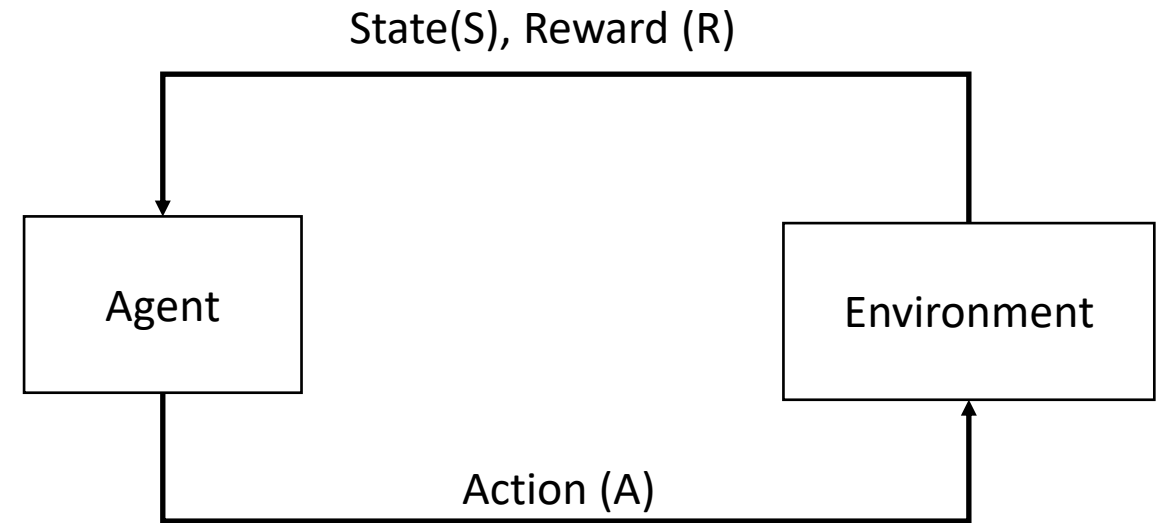
Reinforcement learning environment: No learning (Human agent)

- Goal
 - catch flag
- Environment
 - 4x4 grid world
 - two mines and one flag
- State
 - (x,y) position of an agent
- Reward
 - catch flag: +1
 - step mine: -1
 - otherwise: 0
- Action
 - 0: right, 1: down, 2: left, 3: up
- Episode end
 - Agent reaches one of mines or a flag
 - Every episode restarts from (0,0)




Reinforcement learning: Q-learning (single agent)

- Goal
 - catch flag
- Environment
 - 4x4 grid world
 - two mines and one flag
- State
 - (x,y) position of an agent
- Reward
 - catch flag: +1
 - step mine: -1
 - otherwise: 0
- Action
 - 0: right, 1: down , 2: left, 3: up
- Episode end
 - Agent reaches one of mines or a flag
 - Every episode restarts from (0,0)



Q-learning



Initialize $Q(s, a) = 0, \forall s \in S, a \in A(s)$ ← `agent_init()`

Repeat (for each episode):

Initialize S

Repeat (for each step of episode:)

Choose A from current state S using policy derived from Q (e.g. ϵ -greedy) } `agent_action()`

Take action A

Observe next state S' and R

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$ } `agent_update()`


$S \leftarrow S'$

Until S is terminal

Agent: Action

- Policy: ϵ -Greedy
- Action
0: right, 1: down , 2: left, 3: up
- S_n : current state
- You need to decrease ϵ every episode
ex) initial value: 1.0 -> 0.1 (minimum value)
Exploration is more important from beginning
- Exploitation: Make the best decision given current information
- Exploration: Gather more information

```
if (uniform(0,1) <  $\epsilon$ )  
    a = uniform(0, #actions) ← Exploration  
else  
    a =  $\arg \max_{a'} Q(S_n, a')$  ← Exploitation (Greedy policy)
```

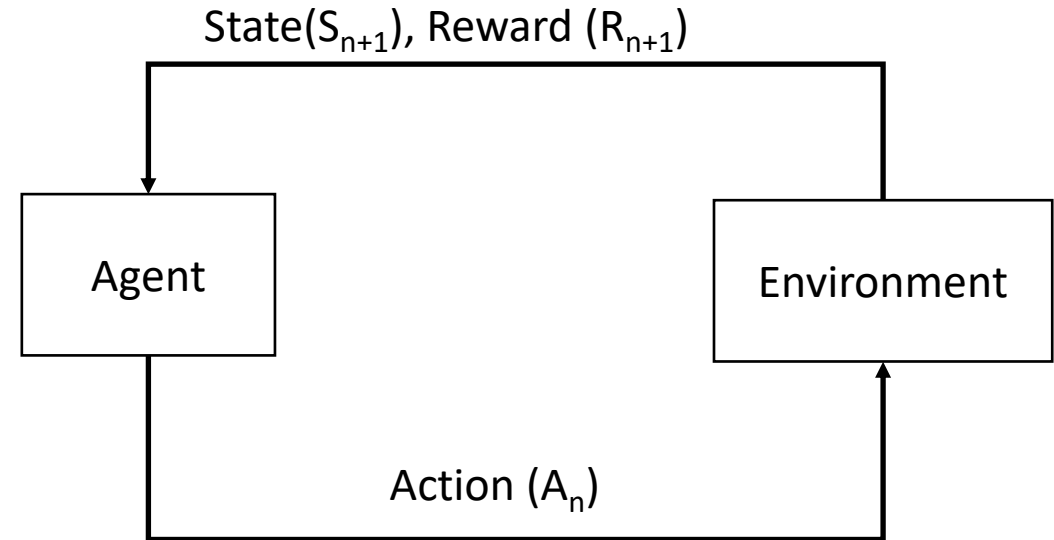


Agent: update Q-table: Q-learning

- Update Q-table
- R_{n+1} : reward
- S_{n+1} : next state
- S_n : current state

α : learning rate



γ : discount factor



$$Q(S_n, A_n) += \overset{\text{Gradient ascent}}{\alpha} \left[\underset{\substack{\text{error} \downarrow \\ \text{prediction to Q} \uparrow}}{R_{n+1} + \gamma \max_{a'} Q(S_{n+1}, a')} - Q(S_n, A_n) \right]$$

Q-table

- This is an example of Q-table
- Action A
0: right, 1: down, 2: left, 3: up



$Q(S_n, A)$, where $S_n = (0, 3)$


$Q(S_n, A)$, where $S_n = (1, 3)$

Q(S,A) table

Optimal Policy: Greedy after Q-learning

- For given state, S_n

$$a = \arg \max_{a'} Q(S_n, a')$$



Source code view

```
if (m_episode == 0 && m_steps==0) {// only for first episode
    env.reset(m_sid);
    agent_init(); // init Q table + self initialization
} else {
    active_agent = checkstatus(board, env.m_state, flag_agent);

    if (m_newepisode) {
        env.reset(m_sid); // clear buffer
        float epsilon = agent_adjustepsilon(); // adjust epsilon
        m_steps = 0;
        printf("EP=%4d, _eps=%4.3f\n", m_episode, epsilon);
        m_episode++;
    } else {
        short* action = agent_action(env.d_state[m_sid]);
        env.step(m_sid, action);
        agent_update(env.d_state[m_sid], env.d_state[m_sid ^ 1], env.d_reward);

        m_sid ^= 1;
        episode = m_episode;
        steps = m_steps;
    }
}
m_steps++; env.render(board, m_sid); return m_newepisode;
```

Algorithm view

Initialize $Q(s, a) = 0, \forall s \in S, a \in A(s)$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

Choose A from current state S using policy derived from Q (e.g. ϵ -greedy)

Take action A

Observe next state S' and R

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$


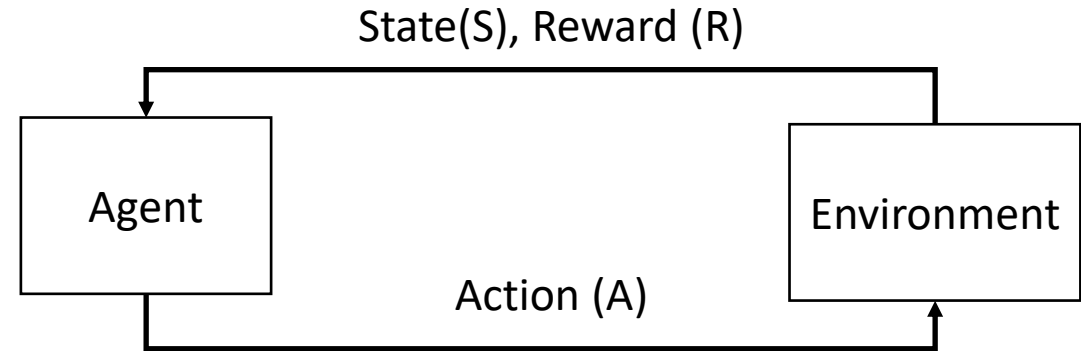
Until S is terminal

Exchange data types

- S_n : current state
env.d_state[m_sid]: *int2 for each agent
- S_{n+1} : next state
env.d_state[m_sid^1]: *int2 for each agent
- Each state indicates of position
d_state[agent_id].x: x position
d_state[agent_id].y: y position
- Reward:
env.d_reward: *float for each agent
d_reward[agent_id]: catch flag: +1, step mine: -1, otherwise: 0

Reinforcement learning: Q-learning

- Current status
(x, y): position of an agent
`env.d_state[m_sid]`
- Update agent
current status: `env.d_state[m_sid]`
next status: `env.d_state[m_sid ^ 1]`
reward: `env.d_reward`
- Action
0: right, 1: bottom, 2: left, 3: top
- Agent status
inactive (nonzero reward) or active (zero reward)



■ curand_uniform usage

- Onetime allocation for every thread

- `curandState *states`
- `cudaMalloc((void **)&states, sizeof(curandState) * threads_per_block * blocks_per_grid)`

- You should pass a pointer of **curandState**

`__global__ void device_api_kernel(curandState *states, float *out, int N)`

- Onetime Initialization of a kernel: All the threads need to maintain their own states

```
__global__ void init_randstate(int size, curandState *state)
{
    int tid = blockIdx.x*blockDim.x + threadIdx.x;
    curand_init(clock() + tid, tid, 0, &state[tid]);
}
```

- `__global__ void kernel_fun (curandState *d_randstate)`
 - `int tid = blockIdx.x*blockDim.x + threadIdx.x;`
 - Just call `curand_uniform(&d_randstate[tid])`
 - **Not necessary to reinitialize curandState**